

ls
mkdir
cat
put

//Commands:

// to check if hadoop is running

jps

->

hadoop@pict-OptiPlex-3020:~\$ jps
18882 NodeManager
19876 Jps
18726 ResourceManager
18486 SecondaryNameNode
18071 NameNode
18265 DataNode

// stopping all the nodes:

stop-all.sh

->

WARNING: Stopping all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: Use CTRL-C to abort.
Stopping namenodes on [localhost]
Stopping datanodes
Stopping secondary namenodes [pict-OptiPlex-3020]
Stopping nodemanagers
Stopping resourcemanager

jps

-> 20662 Jps

start-dfs.sh

-> Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [pict-OptiPlex-3020]

start-yarn.sh

-> Starting resourcemanager
Starting nodemanagers

hdfs dfs -mkdir /practice

hdfs dfs -put /home/hadoop/test.txt /practice

yarn jar /home/hadoop/hadoop-3.3.5/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.3.5.jar wordcount /practice/test.txt /outputpractice

->

2023-05-09 12:44:26,987 INFO client.DefaultNoHARMAFailoverProxyProvider: Connecting to ResourceManager at /127.0.0.1:8032
2023-05-09 12:44:27,306 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1683616252585_0001
2023-05-09 12:44:27,531 INFO input.FileInputFormat: Total input files to process : 1
2023-05-09 12:44:27,664 INFO mapreduce.JobSubmitter: number of splits:1
2023-05-09 12:44:27,822 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1683616252585_0001
2023-05-09 12:44:27,822 INFO mapreduce.JobSubmitter: Executing with tokens: []
2023-05-09 12:44:27,942 INFO conf.Configuration: resource-types.xml not found
2023-05-09 12:44:27,943 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2023-05-09 12:44:28,133 INFO impl.YarnClientImpl: Submitted application application_1683616252585_0001
2023-05-09 12:44:28,183 INFO mapreduce.Job: The url to track the job: http://pict-OptiPlex-3020:8088/proxy/application_1683616252585_0001/
2023-05-09 12:44:28,184 INFO mapreduce.Job: Running job: job_1683616252585_0001
2023-05-09 12:44:34,266 INFO mapreduce.Job: Job job_1683616252585_0001 running in uber mode : false
2023-05-09 12:44:34,268 INFO mapreduce.Job: map 0% reduce 0%
2023-05-09 12:44:38,313 INFO mapreduce.Job: map 100% reduce 0%
2023-05-09 12:44:43,350 INFO mapreduce.Job: map 100% reduce 100%
2023-05-09 12:44:44,368 INFO mapreduce.Job: Job job_1683616252585_0001 completed successfully
2023-05-09 12:44:44,440 INFO mapreduce.Job: Counters: 54

File System Counters

FILE: Number of bytes read=281
FILE: Number of bytes written=552433
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=257
HDFS: Number of bytes written=183
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=1911
Total time spent by all reduces in occupied slots (ms)=2116
Total time spent by all map tasks (ms)=1911
Total time spent by all reduce tasks (ms)=2116
Total vcore-milliseconds taken by all map tasks=1911

Total vcore-milliseconds taken by all reduce tasks=2116
 Total megabyte-milliseconds taken by all map tasks=1956864
 Total megabyte-milliseconds taken by all reduce tasks=2166784
Map-Reduce Framework
 Map input records=2
 Map output records=26
 Map output bytes=257
 Map output materialized bytes=281
 Input split bytes=104
 Combine input records=26
 Combine output records=23
 Reduce input groups=23
 Reduce shuffle bytes=281
 Reduce input records=23
 Reduce output records=23
 Spilled Records=46
 Shuffled Maps =1
 Failed Shuffles=0
 Merged Map outputs=1
 GC time elapsed (ms)=35
 CPU time spent (ms)=970
 Physical memory (bytes) snapshot=488665088
 Virtual memory (bytes) snapshot=5454766080
 Total committed heap usage (bytes)=427819008
 Peak Map Physical memory (bytes)=268034048
 Peak Map Virtual memory (bytes)=2722947072
 Peak Reduce Physical memory (bytes)=220631040
 Peak Reduce Virtual memory (bytes)=2731819008
Shuffle Errors
 BAD_ID=0
 CONNECTION=0
 IO_ERROR=0
 WRONG_LENGTH=0
 WRONG_MAP=0
 WRONG_REDUCE=0
File Input Format Counters
 Bytes Read=153
File Output Format Counters
 Bytes Written=183

hdfs dfs -cat /outputpractice/part-r-00000

```

-> Deploying 1
Hadoop      2
a           2
acquainted  1
and         1
basic       1
commands    1
concepts    1
get         1
  
```

great	1	
hadoop2		
hello	1	
hi	1	
is	1	
node	1	
on	1	
pict	1	
services		1
single	1	
to	1	
way	1	
with	1	
yourself		1